

### Task3

1. Compile both OMP and CUDA versions of your selected benchmarks. Do you need to make any changes in Makefile?

We need to modify the rodinia\_3.1/cuda/lud/cuda/Makefile if we want to run the lud.

```
remove: -arch=sm_13 \  
in NVCCFLAGS += section
```

For rodinia\_3.1/cuda/b+tree/Makefile,

```
remove line CUDA_FLAG = -arch sm_20 and all variables related to CUDA_FLAG
```

2. Ensure the same input problem is used for OMP and CUDA versions. Report and compare their execution time.

```
root@Daniel:~/DD2360/Assignment_1/rodinia_3.1/rodinia_3.1/openmp/lud# bash run  
Generate input matrix internally, size =256  
Creating matrix internally size=256  
running OMP on host  
Time consumed(ms): 27.457000  
  
root@Daniel:~/DD2360/Assignment_1/rodinia_3.1/rodinia_3.1/cuda/lud# bash run  
WG size of kernel = 16 X 16  
Generate input matrix internally, size =256  
Creating matrix internally size=256  
Before LUD  
Time consumed(ms): 3.209000  
After LUD  
>>>Verify<<<<
```

```
root@Daniel:~/DD2360/Assignment_1/rodinia_3.1/rodinia_3.1/openmp/b+tree# bash run
Input File: ../../data/b+tree/mil.txt
Command File: ../../data/b+tree/command.txt
Command Buffer:
j 6000 3000
k 10000

Getting input from file core...
Transforming data to a GPU suitable structure...
Tree transformation took 0.098952
Waiting for command
>
*****command: j count=6000, rSize=6000
Time spent in different stages of CPU/MCPU KERNEL:
0.000005000000 s, 0.047732695937 % : MCPU: SET DEVICE
0.010470000096 s, 99.952270507812 % : CPU/MCPU: KERNEL
Total time:
0.010475000367 s
> > > > > > > > >
*****command: k count=10000
Time spent in different stages of CPU/MCPU KERNEL:
0.000001000000 s, 0.008475294337 % : MCPU: SET DEVICE
0.011797999963 s, 99.991523742676 % : CPU/MCPU: KERNEL
Total time:
0.011799000204 s
> > > > > > > > >
root@Daniel:~/DD2360/Assignment_1/rodinia_3.1/rodinia_3.1/openmp/b+tree#
```

```

root@Daniel:~/DD2360/Assignment_1/rodinia_3.1/rodinia_3.1/cuda/b+tree# bash run
WG size of kernel 1 & 2 = 256
Selecting device 0
Input File: ../../data/b+tree/mil.txt
Command File: ../../data/b+tree/command.txt
Command Buffer:
j 6000 3000
k 10000
Getting input from file ../../data/b+tree/mil.txt...
Transforming data to a GPU suitable structure...
Tree transformation took 0.403506
Waiting for command
>
*****command: j count=6000, rSize=6000
knodes_elem=7874, knodes_unit_mem=2068, knodes_mem=16283432
# of blocks = 6000, # of threads/block = 256 (ensure that device can handle)
Time spent in different stages of GPU_CUDA KERNEL:
0.864696025848 s, 98.529289245605 % : GPU: SET DEVICE / DRIVER INIT
0.002187999897 s, 0.249315470457 % : GPU MEM: ALO
0.008689999580 s, 0.990197181702 % : GPU MEM: COPY IN
0.000292999990 s, 0.033386394382 % : GPU: KERNEL
0.000315000012 s, 0.035893220454 % : GPU MEM: COPY OUT
0.001420999994 s, 0.161918312311 % : GPU MEM: FRE
Total time:
0.877602994442 s
> > > > > > > > >
*****command: k count=10000
records_elem=1000000, records_unit_mem=4, records_mem=4000000
knodes_elem=7874, knodes_unit_mem=2068, knodes_mem=16283432
# of blocks = 10000, # of threads/block = 256 (ensure that device can handle)
Time spent in different stages of GPU_CUDA KERNEL:
0.000015000000 s, 0.278035223484 % : GPU: SET DEVICE / DRIVER INIT
0.001618999988 s, 30.009265899658 % : GPU MEM: ALO
0.002870999975 s, 53.215938568115 % : GPU MEM: COPY IN
0.000209000005 s, 3.873957395554 % : GPU: KERNEL
0.000032000000 s, 0.593141794205 % : GPU MEM: COPY OUT
0.000648999994 s, 12.029656410217 % : GPU MEM: FRE
Total time:
0.005394999869 s
> > > > > > > > >

```

3. Do you observe expected speedup on GPU compared to CPU? Why or Why not?

There's speedup for lud, 27.45ms is reduced to 3.21ms when running on gpu.

For b+tree, 0.877602994442s for gpu and 0.017s for cpu.

Therefore, not all algorithms could benefit from parallelism.